

+ Code + Text

```
#installing kaggle library
! pip install kaggle
```

```
Requirement already satisfied: kaggle in /usr/local/lib/python3.10/dist-packages (1.6.17)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.10/dist-packages (from kaggle) (1.16.0)
Requirement already satisfied: certifi>=2023.7.22 in /usr/local/lib/python3.10/dist-packages (from kaggle) (2024.8.30)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.10/dist-packages (from kaggle) (2.8.2)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from kaggle) (2.32.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from kaggle) (4.66.6)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.10/dist-packages (from kaggle) (8.0.4)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.10/dist-packages (from kaggle) (2.2.3)
Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (from kaggle) (6.2.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from bleach->kaggle) (0.5.1)
Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.10/dist-packages (from python-slugify->kaggle) (1.3)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->kaggle) (3.4.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->kaggle) (3.10)
```

```
#configuring the path of Kaggle.json file
! mkdir ~/.kaggle
! cp kaggle.json ~/.kaggle/
! chmod 600 ~/.kaggle/kaggle.json
```

```
mkdir: cannot create directory '/root/.kaggle': File exists
```

[] Suggested code may be subject to a license | AashaiAvadhani1/ChatGPT-Or-Not | Adityapandey0987/fake_news_detection

```
#importing the dependencies
import numpy as np
import pandas as pd
import re
from nltk.corpus import stopwords
from nltk.stem.porter import PorterStemmer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

```
[ ] import nltk
nltk.download('stopwords')
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
True
```

+ Code + Text

Reconnect + Gemini ↕

```
[ ] #printing the stopwords in English
print(stopwords.words('english'))
```

```
'down', 'in', 'out', 'on', 'off', 'over', 'under', 'again', 'further', 'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any', 'both', 'each', 'few', 'more', 'most', 'other', 'some', 'such', 'no', 'nor', 'not', 'only', 'or'
```

```
[ ] #data processing
#loading the data from csv file to dataframe
twitter_data = pd.read_csv('/content/training.1600000.processed.noemoticon.csv', encoding = 'ISO-8859-1')
```

```
[ ] #checking the number of rows and columns
twitter_data.shape
```

```
(608611, 1)
```

```
[ ] #printing the first 5 rows of the dataframe
twitter_data.head()
```

| | | | | | amazing singer LUV HER |
|---|------------|------------------------------|----------|---------------|--|
| 4 | 1753806888 | Sun May 10 03:33:08 PDT 2009 | NO_QUERY | yayitsezekiel | Going to bed..I love the weekends |
| | 1753806977 | Sun May 10 03:33:10 PDT 2009 | NO_QUERY | AubweeMowee | Finally home and ready for bed! night |
| | 1753806984 | Sun May 10 03:33:10 PDT 2009 | NO_QUERY | rumpleEs | @nicolejacinto Ahh, your baby is so cute! Happ.. |
| | 1753807001 | Sun May 10 03:33:11 PDT 2009 | NO_QUERY | RinoaTakako | @verflucht Thanks |
| | 1753807031 | Sun May 10 03:33:11 PDT 2009 | NO_QUERY | kravmascara | Kids b'day party/picnic in canazarno park... |

```
#naming the column and reading the dataset again
```

```
column_names = ['target', 'id', 'date', 'flag', 'user', 'text']
twitter_data = pd.read_csv('/content/training.1600000.processed.noemoticon.csv', encoding = 'ISO-8859-1', names = column_names)
```

```
!python-input-15-6d8abed2101b-v4: DtypeWarning: Columns (0) have mixed types. Specify dtype option on import or set low_memory=False.
twitter_data = pd.read_csv('/content/training.1600000.processed.noemoticon.csv', encoding = 'ISO-8859-1', names = column_names)
```

```
[ ] #checking the number of rows and columns
twitter_data.shape
```

```
(608612, 6)
```

+ Code + Text

```
[ ] #0--> negative tweet
#1--> positive tweet
#streaming
#streaming is the process of reducing a word to its root word
port_stem = PorterStemmer()

[ ] Suggested code may be subject to a license | Adityapandey0987/fake_news_detection | mikael-kumsa/Twitter-Sentiment-Analysis
def stemming(content):
    stemmed_content = re.sub('[^a-zA-Z]', ' ', content)
    stemmed_content = stemmed_content.lower()
    stemmed_content = stemmed_content.split()
    stemmed_content = [port_stem.stem(word) for word in stemmed_content if not word in stopwords.words('english')]
    stemmed_content = ' '.join(stemmed_content)
    return stemmed_content

[ ] print(twitter_data['text'].head())
print(twitter_data['text'].dtype)

0      None
1      Going to bed...I love the weekends
2      Finally home and ready for bed! night!
3      @nicolejacinto Ahh, your baby is so cute! Happ...
4      @verflucht Thanks
Name: text, dtype: object
object

[ ] twitter_data['text'] = twitter_data['text'].fillna('').astype(str)

[ ] def stemming(text):
    if not isinstance(text, str):
        return ''

twitter_data['stemmed_content'] = twitter_data['text'].apply(stemming)
```

+ Code + Text

```
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, random_state=2)
print(X.shape, X_train.shape, X_test.shape)

(680611,) (544488,) (136123,)

[ ] print(X_train)
[None None None ... None None None]

[ ] print(X_test)
[None None None ... None None None]

[ ] print([i for i, x in enumerate(X_train) if x is None or x == ""]) # Indices of None or empty string in X_train
print([i for i, x in enumerate(X_test) if x is None or x == ""]) # Indices of None or empty string in X_test

[ ] X_train = np.array([x if x is not None and x != "" else "missing" for x in X_train])
X_test = np.array([x if x is not None and x != "" else "missing" for x in X_test])

from sklearn.feature_extraction.text import TfidfVectorizer

vectorizer = TfidfVectorizer()

X_train = vectorizer.fit_transform(X_train)
X_test = vectorizer.transform(X_test)
```

+ Code + Text

```
[ ] print(X)
```

```
[None None None ... None None None]
```

```
[ ] print(Y)
```

```
[' amazing singer LUV HER'' '4' '4' ... 1 1 1]
```

```
[ ] print(Y[:5]) # Display the first 5 elements
```

```
[' amazing singer LUV HER'' '4' '4' '4' '4']
```

```
[ ] print(Y.dtype)
```

```
object
```

```
Y = Y.astype(str)
```

```
[ ] import pandas as pd  
Y = pd.to_numeric(Y, errors='coerce') # Convert; invalid entries become NaN
```

```
[ ] from sklearn.utils import shuffle  
  
valid_indices = ~np.isnan(Y) # Identify non-NaN indices  
X, Y = X[valid_indices], Y[valid_indices] # Filter valid data
```

```
[ ] print(np.unique(Y))
```

```
[1. 4.]
```

+ Code + Text

```
from sklearn.model_selection import train_test_split  
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, random_state=2)  
print(X.shape, X_train.shape, X_test.shape)
```

```
(680611,) (544488,) (136123,)
```

```
[ ] print(X_train)
```

```
[None None None ... None None None]
```

```
[ ] print(X_test)
```

```
[None None None ... None None None]
```


```
[ ] print([i for i, x in enumerate(X_train) if x is None or x == ""]) # Indices of None or empty string in X_train  
print([i for i, x in enumerate(X_test) if x is None or x == ""]) # Indices of None or empty string in X_test
```

```
[ ] X_train = np.array([x if x is not None and x != "" else "missing" for x in X_train])  
X_test = np.array([x if x is not None and x != "" else "missing" for x in X_test])
```


```
from sklearn.feature_extraction.text import TfidfVectorizer  
  
vectorizer = TfidfVectorizer()  
  
X_train = vectorizer.fit_transform(X_train)  
X_test = vectorizer.transform(X_test)
```

+ Code + Text

 print(X_train)

 (0, 0) 1.0
(1, 0) 1.0
(2, 0) 1.0
(3, 0) 1.0
(4, 0) 1.0
(5, 0) 1.0
(6, 0) 1.0
(7, 0) 1.0
(8, 0) 1.0
(9, 0) 1.0
(10, 0) 1.0
(11, 0) 1.0
(12, 0) 1.0
(13, 0) 1.0
(14, 0) 1.0
(15, 0) 1.0
(16, 0) 1.0
(17, 0) 1.0
(18, 0) 1.0
(19, 0) 1.0
(20, 0) 1.0
(21, 0) 1.0
(22, 0) 1.0
(23, 0) 1.0
(24, 0) 1.0
:
(544463, 0) 1.0
(544464, 0) 1.0
(544465, 0) 1.0
(544466, 0) 1.0
(544467, 0) 1.0
(544468, 0) 1.0
(544469, 0) 1.0
(544470, 0) 1.0
(544471, 0) 1.0
(544472, 0) 1.0
(544473, 0) 1.0
(544474, 0) 1.0
(544475, 0) 1.0
(544476, 0) 1.0
(544477, 0) 1.0
(544478, 0) 1.0
(544479, 0) 1.0
(544480, 0) 1.0
(544481, 0) 1.0
(544482, 0) 1.0
(544483, 0) 1.0
(544484, 0) 1.0
(544485, 0) 1.0
(544486, 0) 1.0
(544487, 0) 1.0

 print(X_test)

 (0, 0) 1.0
(1, 0) 1.0
(2, 0) 1.0
(3, 0) 1.0
(4, 0) 1.0
(5, 0) 1.0
(6, 0) 1.0
(7, 0) 1.0
(8, 0) 1.0
(9, 0) 1.0
(10, 0) 1.0
(11, 0) 1.0
(12, 0) 1.0
(13, 0) 1.0
(14, 0) 1.0
(15, 0) 1.0
(16, 0) 1.0
(17, 0) 1.0
(18, 0) 1.0
(19, 0) 1.0
(20, 0) 1.0
(21, 0) 1.0
(22, 0) 1.0
(23, 0) 1.0
(24, 0) 1.0
:
(136098, 0) 1.0
(136099, 0) 1.0
(136100, 0) 1.0
(136101, 0) 1.0
(136102, 0) 1.0
(136103, 0) 1.0
(136104, 0) 1.0
(136105, 0) 1.0
(136106, 0) 1.0
(136107, 0) 1.0
(136108, 0) 1.0
(136109, 0) 1.0
(136110, 0) 1.0
(136111, 0) 1.0
(136112, 0) 1.0
(136113, 0) 1.0
(136114, 0) 1.0
(136115, 0) 1.0
(136116, 0) 1.0
(136117, 0) 1.0
(136118, 0) 1.0
(136119, 0) 1.0
(136120, 0) 1.0
(136121, 0) 1.0
(136122, 0) 1.0

+ Code + Text

```
[ ] #training the machine learning model
#logistic Regression
model = LogisticRegression(max_iter=1000)
```

```
[ ] model.fit(X_train, Y_train)
```

```
LogisticRegression
LogisticRegression(max_iter=1000)
```

```
[ ] #Model Evaluation
#Accuracy Score
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(Y_train, X_train_prediction)
```

```
[ ] print('Accuracy score on the training data:', training_data_accuracy)
```

```
Accuracy score on the training data: 0.8074209165307592
```

```
[ ] X_test_prediction = model.predict(X_test)
training_data_accuracy = accuracy_score(Y_test, X_test_prediction)
```

```
from sklearn.metrics import accuracy_score

# Assuming your model is named 'model' and you're predicting on 'X_train'
y_pred_train = model.predict(X_train)
training_data_accuracy = accuracy_score(Y_train, y_pred_train)

print('Accuracy score on the training data:', training_data_accuracy)
```

```
Accuracy score on the training data: 0.8074209165307592
```

```
[ ] y_pred_test = model.predict(X_test)
testing_data_accuracy = accuracy_score(Y_test, y_pred_test)

print('Accuracy score on the test data:', testing_data_accuracy)
```

```
Accuracy score on the test data: 0.8074241678481962
```

+ Code + Text

```
[ ] #model accuracy = 77.8%
#saving the trained model
import pickle
filename = 'trained_model.sav'
pickle.dump(model, open(filename, 'wb'))
```

```
[ ] #loading the saved model
loaded_model = pickle.load(open('/content/trained_model.sav', 'rb'))
```

```
[ ] X_new = X_test[200]
print(Y_test[200])

prediction = loaded_model.predict(X_new)
print(prediction)

if(prediction[0]==0):
    print("Negative Tweet")
else:
    print("Positive Tweet")
```

```
4.0
[1.]
Positive Tweet
```

```
X_new = X_test[3]
print(Y_test[3])

prediction = loaded_model.predict(X_new)
print(prediction)

if(prediction[0]==0):
    print("Negative Tweet")
else:
    print("Positive Tweet")
```

```
1.0
[1.]
Positive Tweet
```